
Flash-SD-KDE: Accelerating SD-KDE with Tensor Cores

Anonymous Author(s)

Affiliation

Address

email

Abstract

Score-debiased kernel density estimation (SD-KDE) achieves improved asymptotic convergence rates over classical KDE, but its use of an empirical score has made it significantly slower in practice. We show that by re-ordering the SD-KDE computation to expose matrix-multiplication structure, Tensor Cores can be used to accelerate the GPU implementation. On a 32k-sample 16-dimensional problem, our approach runs up to $47\times$ faster than a strong SD-KDE GPU baseline and $3,300\times$ faster than scikit-learn’s KDE. On a larger 1M-sample 16-dimensional task evaluated on 131k queries, Flash-SD-KDE completes in 2.3 s on a single GPU, making score-debiased density estimation practical at previously infeasible scales. Code available in supplementary material.

1 Introduction

Kernel density estimation (KDE) is a foundational nonparametric tool [Rosenblatt, 1956, Parzen, 1962, Silverman, 1986, Wand and Jones, 1995] that has become a recurring primitive in modern machine learning pipelines. KDE-style density and score estimates underpin anomaly and out-of-distribution detection [Luo and Shrivastava, 2018, Coleman and Shrivastava, 2020], embedding-space dataset comparison and divergence estimation, data-mixture and selection analysis for pre-training and instruction tuning, density-based regularization and sampling, and softmax-style attention [Zandieh et al., 2023, Kacham et al., 2024]. At the scale of these applications, however, two classical limitations become decisive: an $O(h^2)$ leading-bias term that dominates error in moderate dimensions, and an $O(n^2)$ exact evaluation cost — both of which have driven decades of follow-up work.

A long line of work attacks the bias limitation by modifying the kernel or the bandwidth so that the leading bias term cancels. Higher-order kernel constructions [Fan and Hu, 1992, Jones and Signorini, 1997], variable-bandwidth schemes [Abramson, 1982], and diffusion-based estimators [Botev et al., 2010] are all members of this family. Score-debiased kernel density estimation (SD-KDE) [Epstein et al., 2025] is a recent addition: it uses an empirical score $\hat{s}$ to form debiased samples $x_i^{\text{SD}} = x_i + \frac{h^2}{2} \hat{s}(x_i)$, and then evaluates a standard Gaussian KDE on these shifted samples instead of the originals, which improves statistical efficiency over standard KDE while retaining a simple, nonparametric form and preserving non-negativity. In particular, SD-KDE achieves an asymptotic mean integrated squared error (AMISE) of order $O(n^{-8/(d+8)})$, improving the convergence exponent relative to the $O(n^{-4/(d+4)})$ rate of classical KDE tuned with Silverman’s rule of thumb [Silverman, 1986, Wand and Jones, 1995].

These statistical gains come with a significant computational drawback. The empirical score used for debiasing is itself a kernel sum, so it introduces an additional $O(n^2)$ pass over the data, and a naive implementation has roughly the same quadratic cost as KDE itself but with a substantially larger constant factor. Empirically, the score pass dominates end-to-end runtime, accounting for roughly 90% of the total cost in our setting, which has so far made SD-KDE impractical at the scales where

downstream applications operate. Although GPU acceleration is the natural route, the score pass does not match the GEMM shape that modern accelerators are built to execute, and the surrounding scalar operations (norms, exponentials, and atomic accumulations) limit the benefit of generic deep-learning frameworks.

Our central observation is that the quadratic interaction is not, by itself, the obstacle; its dataflow is. Re-ordering the SD-KDE computation exposes hidden matrix-multiplication structure inside both the empirical-score pass and the KDE evaluation. Once that structure is made explicit, the heavy work of SD-KDE is exactly the work that Tensor Cores on modern GPUs are designed to accelerate, and the rest of the implementation reduces to streaming the matrix tiles and managing the surrounding scalar work against a mixed-precision performance budget. The estimator remains exact, but the computation is organized in the hardware-native order. This is what enables SD-KDE to run at scales that were previously infeasible on a single GPU. Beyond the headline synthetic benchmark, we also stress-test the implementation against stronger exact baselines, across a second GPU architecture, under varying query batch sizes, through streaming and Tensor-Core ablations, and on real anomaly-detection benchmarks.

Contributions. We make SD-KDE practical at large scale through three contributions: a GEMM re-ordering that exposes Tensor-Core structure in both the score and the KDE passes, a streaming Triton implementation that avoids materializing pairwise kernel matrices, and a fused first-order surrogate (Flash-Laplace-KDE) that retains the leading bias correction without the explicit score pass. On a 32k-sample 16-D problem, Flash-SD-KDE runs roughly $47.2\times$ faster than a strong PyTorch SD-KDE baseline, $9.0\times$ faster than PyKeOps, and over three orders of magnitude faster than scikit-learn KDE. On a 1M-sample 16-D task evaluated on 131k queries, it completes in 2.3 seconds on a single GPU.

Paper overview. Section 2 surveys related work on KDE, score estimation, fast kernel evaluation, and GPU programming. Section 3 reviews the GPU hardware features that drive the design. Section 4 derives the matrix-form SD-KDE computation and the streaming Tensor-Core kernel. Section 5 introduces the Laplace-corrected surrogate. Section 6 reports the empirical study, including performance optimizations and the oracle-density benchmark, and Section 7 concludes. Additional experiments, including exact-baseline sweeps, cross-GPU results, query-batch sensitivity, ablations, and anomaly-detection benchmarks, appear in Appendix A; a more complete literature review appears in the appendix.

2 Related Work

Kernel density estimation. Kernel density estimation (KDE) is a classical nonparametric approach to density estimation that dates back to the foundational work of Rosenblatt [1956] and Parzen [1962]; modern treatments appear in standard references such as Silverman [1986] and Wand and Jones [1995]. For smooth densities, KDE attains optimal nonparametric rates under appropriate bandwidth scaling, but in practice performance is highly sensitive to bandwidth selection. Rules-of-thumb and plug-in procedures are widely used for their simplicity, but they are typically tuned to unimodal, near-Gaussian targets and can misbehave for multimodal or heavy-tailed distributions [Silverman, 1986, Botev et al., 2010]. This sensitivity is particularly pronounced in moderate and high dimensions, where the bandwidth controls both statistical bias and the computational cost of evaluating kernels at scale.

Bias reduction and adaptive smoothing. A large body of work improves KDE accuracy by reducing leading-order bias or adapting the effective smoothing scale across the sample space. One line of work constructs higher-order bias corrections by modifying kernels or combining estimators so that the dominant $O(h^2)$ bias term cancels, yielding $O(h^4)$ bias under additional smoothness assumptions [Fan and Hu, 1992, Jones and Signorini, 1997]. Another class of methods uses variable bandwidths that shrink in regions of high density and expand in the tails, improving adaptivity without committing to a parametric form [Abramson, 1982, Silverman, 1986]. Complementary to these ideas, diffusion-based estimators view density estimation through the lens of smoothing by a linear diffusion process, providing both adaptivity and a data-driven bandwidth selection mechanism [Botev et al., 2010]. Score-debiased KDE (SD-KDE) [Epstein et al., 2025] fits into this landscape by using score information to correct the leading bias term while retaining a simple kernel estimator structure. Our Laplace-corrected formulation can be seen as an explicit higher-order correction that removes the

91 need to form an empirical score at evaluation time, trading additional kernel evaluations for improved
92 numerical and computational behavior.

93 **Score estimation and score-based modeling.** The score function $\nabla_x \log p(x)$ is a central object in
94 statistics and machine learning. Score matching [Hyvärinen, 2005] provides a way to fit unnormalized
95 models by matching scores rather than likelihoods, and denoising score matching connects score
96 estimation to learning denoisers [Vincent, 2011]. More recently, score-based generative models
97 learn scores of noise-perturbed data distributions with neural networks and sample via Langevin
98 or SDE-based dynamics [Song and Ermon, 2019, Song et al., 2021, Ho et al., 2020]. These works
99 typically target high-dimensional perceptual data and emphasize sampling or likelihood computation,
100 whereas SD-KDE uses an explicit nonparametric density estimator and an analytically-defined score
101 to improve pointwise density estimation. Related to our setting, Stein-based score estimators recover
102 score information from samples via identities that avoid explicit density evaluation [Li and Turner,
103 2017]. Our focus is complementary: we keep the statistical estimator fixed (SD-KDE and its Laplace
104 variant) and address the practical bottleneck that arises when score-corrected estimators are deployed
105 at large scale.

106 **Fast evaluation of kernel sums and density models.** The computational cost of KDE and related
107 kernel methods is dominated by evaluating large Gaussian sums or kernel matrices. Algorithmic
108 accelerations include multipole and series-expansion methods such as the Fast Gauss Transform
109 (FGT) [Greengard and Strain, 1991] and improved variants tailored to Gaussian kernels [Yang et al.,
110 2003]. A different set of approaches uses randomized low-rank approximations (e.g., random Fourier
111 features) to trade bias for computational efficiency [Rahimi and Recht, 2007]. These methods reduce
112 asymptotic complexity but can be sensitive to dimensionality, bandwidth, and target accuracy. In
113 contrast, our work targets the exact SD-KDE computation and focuses on constant-factor reductions
114 that make the estimator practical on modern accelerators.

115 **GPU acceleration and tensor-core programming.** A growing ecosystem of libraries and DSLs
116 targets large kernel computations on GPUs. KeOps provides a memory-efficient reduction framework
117 for kernel and distance matrices and is a strong baseline for Gaussian kernel reductions at scale
118 [Charlier et al., 2021]. General deep learning frameworks such as PyTorch [Paszke et al., 2019] make
119 it easy to express quadratic kernel computations, but their performance depends critically on whether
120 the computation can be lowered to high-throughput primitives. Domain-specific languages such as
121 Triton [Tillet et al., 2019] enable writing custom GPU kernels while still benefiting from compiler
122 support for tiling and scheduling. At the hardware level, modern GPUs provide specialized matrix-
123 multiply units (Tensor Cores) and mixed-precision execution paths that can offer large throughput
124 improvements when computations are expressed in GEMM-like form [Micikevicius et al., 2017,
125 Williams et al., 2009]. Recent “flash”-style algorithms demonstrate that reordering computations to
126 avoid materializing large intermediate matrices and to match the GPU memory hierarchy can yield
127 substantial speedups for quadratic primitives [Dao et al., 2022]. Our contribution follows this direction
128 for score-corrected density estimation: we reorder SD-KDE to expose matrix-multiplication structure
129 and implement a streaming accumulation strategy so that Tensor Cores can be used effectively without
130 incurring prohibitive memory traffic.

131 3 Hardware

132 All experiments run on a workstation equipped with an NVIDIA RTX A6000 (GA102). This GPU
133 exposes 84 streaming multiprocessors (SMs); each SM contains 128 FP32 ALUs and 16 special
134 function units (SFUs). Because the ratio of FP32 ALUs to SFUs is 128:16, we treat one exp (issued
135 on an SFU) as costing the equivalent of $128/16 = 8$ FP32 flops in our models. The card delivers
136 roughly 40 TFLOP/s of peak FP32 throughput and approximately 770 GB/s of GDDR6 bandwidth.

137 The host CPU is a dual-socket AMD EPYC 7763 system (“Milan”) configured with 2×64 cores
138 and two hardware threads per core (256 logical CPUs reported by `lscpu`). The processors boost
139 up to 3.53 GHz, expose 512 MiB of shared L3 cache across the sockets, and provide modern ISA
140 extensions (AVX/AVX2, FMA, BMI1/2, SHA, VAES, etc.).




Figure 1: Runtime comparison for 16-D KDE/SD-KDE across n_{train} up to 32,768 ($n_{\text{test}} = n_{\text{train}}/8$).

4 Method

We begin by recalling the standard KDE and SD-KDE definitions that the rest of the section will accelerate. Given samples $x_i \in \mathbb{R}^d$ and a bandwidth h , a Gaussian KDE is

$$\hat{p}(x) = \frac{1}{nh^d} \sum_{i=1}^n \varphi\left(\frac{x - x_i}{h}\right),$$

and SD-KDE forms debiased samples $x_i^{\text{SD}} = x_i + \frac{h^2}{2} \hat{s}(x_i)$ where $\hat{s}$ is an estimate of the score. In this work we always use an empirical score computed from the KDE itself. Assuming a standard Gaussian kernel φ and writing the KDE explicitly in terms of the samples, the empirical score is

$$\hat{s}(x) = \frac{\nabla_x \hat{p}(x)}{\hat{p}(x)} = \frac{\sum_{i=1}^n [-(x - x_i) \varphi(\frac{x - x_i}{h})]}{h^2 \sum_{i=1}^n \varphi(\frac{x - x_i}{h})}.$$

The score therefore has the same $O(n^2)$ pairwise structure as the KDE sum itself, which is the cost we target in what follows.

Our goal is to take n_{train} samples in d dimensions, form the empirical SD-KDE (score + shift), and evaluate the resulting density on n_{test} queries. Writing the computation in matrix form highlights the parallel structure. Let $x_i \in \mathbb{R}^d$ denote training points and $y_j \in \mathbb{R}^d$ denote query points, with bandwidth h and squared Euclidean distance

$$\|x_i - y_j\|^2 = \|x_i\|^2 + \|y_j\|^2 - 2x_i^\top y_j.$$

Stacking the training samples into $X \in \mathbb{R}^{n_{\text{train}} \times d}$, the empirical score computation is dominated by the train–train Gram matrix

$$G_{\text{score}} = XX^\top \in \mathbb{R}^{n_{\text{train}} \times n_{\text{train}}}.$$

Stacking queries into $Y \in \mathbb{R}^{n_{\text{test}} \times d}$ and denoting the debiased training samples by

$$X^{\text{SD}} = X + \frac{h^2}{2} \hat{s}(X),$$

the final SD-KDE evaluation uses the train–query Gram matrix

$$G_{\text{KDE}} = X^{\text{SD}}Y^\top \in \mathbb{R}^{n_{\text{train}} \times n_{\text{test}}}.$$

On modern NVIDIA GPUs this GEMM can be mapped to Tensor Cores [Micikevicius et al., 2017] and evaluated at 5–10 $\times$ the throughput of standard FP32 SIMT arithmetic. The remaining operations—vector norms, broadcasted additions, and exponentials—are all $O(n_{\text{train}}n_{\text{test}})$ but quickly become a small fraction of the total FLOPs as d grows.

161 The numerator in the empirical score inherits a similar GEMM structure. Its numerator involves
 162 terms of the form

$$\sum_j -(x_i - x_j) \varphi_{ij}, \quad \varphi_{ij} = \exp\left(-\frac{\|x_i - x_j\|^2}{2h^2}\right),$$

163 which naively suggests $O(n_{\text{train}} n_{\text{test}} d)$ additional elementwise arithmetic. Using the identity

$$\sum_j (x_i - x_j) \varphi_{ij} = x_i \sum_j \varphi_{ij} - \sum_j \varphi_{ij} x_j,$$

164 we can decompose the first term in the numerator into a elementwise sum, and the second term in the
 165 numerator into a second operation:

$$T = \Phi X.$$

166 Thus both the KDE evaluation and the SD-KDE score numerator reduce to Tensor-Core-accelerable
 167 matrix multiplies, plus $O(n^2)$ scalar work for the norms and exponentials.

168 4.1 Arithmetic intensity in d dimensions

169 To understand when the high-dimensional SD-KDE becomes compute-bound, we estimate FLOPs,
 170 bytes moved, and arithmetic intensity for the d -dimensional case. Let $n_{\text{train}} = k$ and set $n_{\text{test}} = k/8$
 171 as in the experiments.

172 **Total FLOPs.** The d -dimensional implementation consists of three main matrix-multiply stages:

- 173 1. Score Gram matrix $G_{\text{score}} = XX^\top$: $2dk^2$ FLOPs.
- 174 2. Score numerator $T = \Phi X$: $2dk^2$ FLOPs, plus $4k^2$ scalar FLOPs for norms and distance
 175 terms and $8k^2$ FLOPs for exponentials (counting each exp as 8 FLOPs due to the SFU/FP32
 176 ratio).
- 177 3. Final KDE Gram matrix on debiased data: $2dk(k/8)$ FLOPs, plus $4k(k/8)$ scalar FLOPs
 178 and $8k(k/8)$ FLOPs for exponentials.

179 Aggregating these terms yields

$$\begin{aligned} \text{FLOPs}_d(k) &\approx 4dk^2 + 12k^2 + 2d\frac{k^2}{8} + 12\frac{k^2}{8} \\ &= \left(4d + 12 + \frac{d}{4} + \frac{3}{2}\right)k^2. \end{aligned}$$

180 Specifically for the $d = 16$ dimensional case, substituting $d = 16$ gives

$$\text{FLOPs}_{16}(k) \approx 81.5 k^2,$$

181 which is on the order of 10^{11} FLOPs for $k = 32k$.

182 **Bytes moved.** To better capture implementation details, we count bytes per tile using the launch pa-
 183 rameters that delivered the best runtime ($BLOCK_M = 64$, $BLOCK_N = 1024$). Each tile loads
 184 $64 \times d$ query values once ($4 BLOCK_M d$ bytes), streams $1024 \times d$ training values ($4 BLOCK_N d$
 185 bytes), and writes the partial PDF and weighted sums ($4 BLOCK_M$ and $4 BLOCK_M d$ bytes).
 186 All of this traffic goes to and from GDDR6, so for $d = 16$ a tile moves roughly

$$\begin{aligned} \text{Bytes}_{\text{tile}} &\approx 4(BLOCK_M d + BLOCK_N d \\ &\quad + BLOCK_M + BLOCK_M d) \\ &= 4(2 BLOCK_M d + BLOCK_N d \\ &\quad + BLOCK_M) \\ &\approx 7.4 \times 10^4 \text{ bytes.} \end{aligned}$$

187 The full problem processes $(k/BLOCK_M) \times (k/BLOCK_N)$ such tiles, so the total GDDR6
 188 traffic is

$$\begin{aligned} \text{Bytes}_{16}(k) &\approx \text{Bytes}_{\text{tile}} \times \frac{k}{BLOCK_M} \times \frac{k}{BLOCK_N} \\ &\approx 1.13 k^2 \text{ bytes.} \end{aligned}$$




Figure 2: Oracle error on a 16D mixture-of-Gaussians benchmark. We report MISE and MIAE versus n_{train} for KDE, Flash-Laplace-KDE (fused Laplace correction), non-fused Laplace correction, and Flash-SD-KDE. The Laplace-corrected estimators can be slightly negative, so error is computed in a signed density manner.

189 **Arithmetic intensity.** Dividing FLOPs by bytes, and substituting $d = 16$, gives

$$I_{16}(k) \approx \frac{\text{FLOPs}_{16}(k)}{\text{Bytes}_{16}(k)} \approx \frac{81.5 k^2}{1.13 k^2} \approx 72 \text{ flops/byte}.$$

190 This tile-aware estimate more closely reflects the measured arithmetic intensity. Comparing against
 191 the A6000 specs (Tensor Core peak of 155 TFLOP/s versus ≈ 770 GB/s of memory bandwidth),
 192 the machine balance is roughly 200 flops/byte [Williams et al., 2009]. Since our kernel sustains
 193 over 70 flops/byte, it lies well into the compute-bound regime on the RTX A6000. Nsight Compute
 194 reports an empirical arithmetic intensity of roughly 95 FLOPs/byte for the score kernel on the A6000,
 195 which is broadly in line with the 72 FLOPs/byte predicted by the tile-level model above. Minor
 196 differences stem from additional bookkeeping work (atomics, reductions) and from Nsight’s finer-
 197 grained accounting of texture/cache traffic, but both figures agree that the kernel operates far above
 198 the bandwidth roofline. Because the kernel mixes Tensor-Core GEMMs with FP32 scalar work
 199 (norms, exponentials, atomics), the effective machine balance sits between the tensor-core roof
 200 (≈ 200 flops/byte) and the FP32 roof (≈ 50 flops/byte). The observed 70–95 flops/byte therefore
 201 straddle these two limits: the GEMM portion is partially bandwidth-limited relative to Tensor Cores,
 202 while the scalar portion is compute-limited relative to FP32 ALUs.

203 5 Laplace-corrected KDE

204 Score-debiased KDE offers better asymptotic error rates, but its empirical score step adds a full
 205 quadratic pass over the data. A natural approximation is to replace the explicit debiasing with a
 206 correction to the kernel itself, in the spirit of classical higher-order bias corrections for KDE¹ [Fan
 207 and Hu, 1992, Jones and Signorini, 1997]. This yields a Laplace-corrected KDE that captures the
 208 same leading-order bias reduction of SD-KDE [Epstein et al., 2025] while avoiding the empirical
 209 score computation at the cost of admitting possibly negative values.

210 Let $K_h(x)$ denote the isotropic Gaussian kernel in d dimensions.

$$K_h(x) = \frac{1}{(2\pi)^{d/2} h^d} \exp\left(-\frac{\|x\|^2}{2h^2}\right),$$

211 Let $\hat{p}(x) = \frac{1}{n} \sum_i K_h(x - x_i)$ be the corresponding KDE. The Laplace-corrected kernel is defined by

$$K_h^{\text{LC}}(x) = K_h(x) - \frac{h^2}{2} \Delta K_h(x).$$

¹Laplace kernel has 0 second moment, but non-zero fourth moment, so it is a fourth-order kernel.

212 For the Gaussian kernel, the Laplacian has a closed form.

$$K_h^{\text{LC}}(x) = K_h(x) \left(1 + \frac{d}{2} - \frac{\|x\|^2}{2h^2} \right),$$

213 The estimator becomes

$$\hat{p}_{\text{LC}}(x) = \frac{1}{n} \sum_{i=1}^n K_h^{\text{LC}}(x - x_i).$$

214 **Connection to SD-KDE.** SD-KDE [Epstein et al., 2025] forms debiased samples $x_i^{\text{SD}} = x_i +$
 215 $\frac{h^2}{2} \hat{s}(x_i)$ using an empirical score $\hat{s}$. Evaluating a KDE on the shifted samples and linearizing in
 216 h^2 yields a correction proportional to ΔK_h . When $\hat{s}$ is taken to be the KDE score, this linearized
 217 estimator coincides with the Laplace-corrected kernel above. Thus $\hat{p}_{\text{LC}}$ can be viewed as a first-order
 218 approximation to SD-KDE that preserves its leading bias reduction while avoiding a score pass. The
 219 heat-semigroup derivation showing the shared $O(h^4)$ leading-bias behavior of Laplace-corrected
 220 KDE and oracle SD-KDE, together with the caveat for empirical scores, is given in Appendix B.2.

221 **Why we consider this correction.** The Laplace-corrected estimator is appealing when the full
 222 SD-KDE score is too expensive or when a fast bias-reduction surrogate is sufficient. It requires
 223 only the same pairwise distances and exponentials as standard KDE, plus a simple affine factor in
 224 the kernel value. The correction can be negative for large $\|x - x_i\|$, so $\hat{p}_{\text{LC}}$ is not guaranteed to be
 225 nonnegative.

226 **Kernel fusion opportunity.** Computing K_h^{LC} reuses the same scaled distances used to compute
 227 K_h . A fused kernel can compute $\phi = \exp(-\frac{1}{2} \text{scaled})$ and apply the Laplace factor inside the same
 228 pass. A non-fused implementation must either recompute distances or materialize intermediates in
 229 a second kernel. This increases memory traffic and kernel launches. We therefore treat the fused
 230 Laplace-corrected kernel as a distinct fast path and refer to it as *Flash-Laplace-KDE*.

231 6 Results

232 We now summarize the empirical behavior of the 16-D implementation on the RTX A6000. For each
 233 n_{train} in $\{2\text{k}, 4\text{k}, 8\text{k}, 16\text{k}, 32\text{k}\}$ we set $n_{\text{test}} = n_{\text{train}}/8$, draw data from a simple 16-D Gaussian mixture,
 234 and run three baselines: scikit-learn KDE [Pedregosa et al., 2011], Torch SD-KDE (GEMM-based,
 235 implemented in PyTorch [Paszke et al., 2019]), and Flash-SD-KDE (Tensor-Core GEMMs for both
 236 KDE and score, implemented in Triton [Tillet et al., 2019]). Figure 1 shows that the Flash-SD-KDE
 237 rapidly outpaces both baselines as n grows. Appendix A.1 expands this comparison to a fuller set of
 238 exact baselines, adding `torch.compile` and PyKeOps across the full n_{train} sweep. Flash-SD-KDE
 239 is the fastest exact method at every measured problem size. In the averaged $n_{\text{train}} = 32,768$ setting, it
 240 is roughly $9.0\times$ faster than PyKeOps, $14.3\times$ faster than `torch.compile`, $47.2\times$ faster than eager
 241 PyTorch SD-KDE, and over three orders of magnitude faster than scikit-learn. The ordering is not
 242 specific to the RTX A6000: Appendix A.2 repeats the same 16-D sweep on a Tesla V100-PCIE-16GB,
 243 where Flash-SD-KDE again remains the fastest method at every measured size.

244 We also measure utilization by combining the flop model above with the measured runtimes. As
 245 shown in Figure 3, the Flash-SD-KDE GPU utilization is high into the multi-digit range once n_{train}
 246 exceeds 8k, despite a lot of operations such as exponential computations that can’t utilize tensor
 247 cores. This confirms that the 16-D Tensor-Core formulation is firmly compute-bound and that
 248 additional tuning effort should focus on kernel fusion and occupancy rather than memory traffic.
 249 The score-dominated profile also makes the end-to-end runtime nearly flat in the query batch size
 250 over a wide practical range. Appendix A.3 fixes $n_{\text{train}} = 32,768$ and sweeps n_{test} from 4 to 16,384;
 251 Flash-SD-KDE remains between 1.91 and 2.75 ms, while the baselines show a larger increase at the
 252 largest query batches. This regime is common in anomaly and out-of-distribution scoring, where
 253 many workflows compare a small set of candidate points against a large reference sample.

254 6.1 Oracle density benchmark for Laplace correction

255 Figure 2 reports oracle accuracy on the 16-D mixture benchmark as a function of n_{train} , using Mean
 256 Integrated Squared Error (MISE) on the left and Mean Integrated Absolute Error (MIAE) on the




Figure 3: Utilization (percentage of RTX A6000 Tensor Core peak) for the 16-D SD-KDE pipeline, computed via the flop model from Section 4; bars are annotated with the observed runtimes (ms).

right. Across the sweep, Flash-SD-KDE substantially improves on vanilla KDE, indicating that the baseline KDE estimator is a poor fit in this regime. The Laplace-corrected variants achieve the lowest MISE throughout, and the fused Flash-Laplace-KDE curve overlaps the non-fused implementation, showing that fusion is an implementation optimization rather than a change in the estimator. In MIAE, Flash-SD-KDE attains the lowest error, while the Laplace-corrected estimators show a sharp drop between the smallest n_{train} and 10^4 , followed by a small non-monotonic “kink” (a slight increase at intermediate n_{train}). The kink lies within the uncertainty bands and is consistent with finite-sample variability and the greater sensitivity of an absolute-error metric to the signed tail correction. Both metrics continue to decrease as n_{train} grows. Because the Laplace-corrected kernel is not guaranteed to remain nonnegative, we compute MISE/MIAE on the signed estimator and separately log the integrated negative mass as a diagnostic. Appendix B.2 reports the corresponding 1D oracle-density and fusion-runtime experiments. The 1D results show the same estimator-level behavior: Flash-Laplace-KDE matches the non-fused Laplace correction in error, but is consistently faster because the correction is applied inside the same kernel pass.

6.2 Real-data anomaly detection

To test whether the accelerated estimator remains useful outside synthetic mixtures, Appendix A.7 evaluates Flash-SD-KDE as an inverse-density anomaly score on real anomaly-detection benchmarks from the ACE study. On Statlog Shuttle, Flash-SD-KDE recovers 807 of 879 true outliers at the chosen threshold, compared with 610 for the strongest published non-Flash baseline in the ACE table and 273 for ACE itself, while running in 0.60 s on the A6000. On ALOI, it recovers 437 true outliers, within 20% of the best published recovery, while running in 0.09 s on the A6000. On a larger public KDD-Cup 99 HTTP release with $n = 620,098$ and $d = 29$, Flash-SD-KDE achieves ROC-AUC 0.92 in 7.27 s. Full thresholding details and baseline tables are reported in Appendix A.7.

6.3 Performance optimizations

Nsight Systems traces show that roughly 95% of the SD-KDE runtime at $n_{\text{train}} = 32\text{k}$, $n_{\text{test}} = 4\text{k}$ is spent computing the empirical score. We therefore tune the score-kernel launch configuration; the full grid is given in Appendix A.4. The best setting uses $BLOCK_M = 64$, $BLOCK_N = 1024$, $num_warps = 2$, and $num_stages = 2$, increasing measured FLOP utilization by more than $2\times$ relative to smaller-tile settings.

Beyond the launch-parameter sweep, two design choices account for most of the speedup over a standard tiling implementation:

1. **Tensor-Core tiling:** all high-dimensional dot products are processed as 16×16 tiles with Tensor Core acceleration, allowing the GEMM portion of the score and KDE kernels to run near the TF32 roofline.
2. **Streaming accumulation:** we never materialize full $n_{\text{train}} \times n_{\text{train}}$ (or query) matrices; instead we stream tiles through registers and rely on atomic reductions to keep global-memory traffic linear in n .

Appendices A.5 and A.6 quantify the factors. At $n_{\text{train}} = 65,536$ and $n_{\text{test}} = 8,192$, streaming reduces peak extra GPU allocation from ~ 16 GB to ~ 16 MB and gives a $55.7\times$ speedup over full materialization. Holding the streamed implementation fixed, enabling Tensor-Core tiling gives a $4.94\times$ speedup at $n_{\text{train}} = 32,768$ and remains near $5\times$ at the largest sizes. Thus streaming makes the exact quadratic computation fit in memory, while Tensor Cores supply the large- n throughput.

Together these optimizations push the kernels close to the hardware limit. Nsight Compute’s “Speed of Light” report shows $\sim 68\%$ SM throughput and roughly 90% L1 throughput for the score kernel at $n_{\text{train}} = 32\text{k}$ and $n_{\text{test}} = 4\text{k}$, which is consistent with the tensor-core utilization trends in Figure 3: we cannot reach the nominal Tensor Core peak because the kernel still executes substantial non-Tensor-Core work (norms, exponentials, atomics). At these utilization levels we stopped further low-level tuning, since additional optimizations would likely yield only modest incremental gains.

7 Conclusion

Score-debiased kernel density estimation (SD-KDE) offers improved asymptotic accuracy over classical KDE, but its reliance on an empirical score has historically made it substantially slower in practice. This paper shows that the dominant cost of empirical SD-KDE is not inherently prohibitive: by re-ordering the computation to expose matrix-multiplication structure, both KDE evaluation and the score numerator can be expressed as Tensor Core-acceleratable GEMMs plus lightweight elementwise operations, while avoiding materializing full pairwise interaction matrices via streaming accumulation. In the 16-D setting, this hardware-aligned formulation yields large end-to-end gains—up to $47\times$ faster than a strong Torch SD-KDE baseline and $3,300\times$ faster than scikit-learn KDE—and scales to previously infeasible regimes, completing SD-KDE on $\sim 1\text{M}$ training points and $\sim 131\text{k}$ queries in 2.3 s on a single GPU.

Overall, our results suggest that practical nonparametric estimators can often be unlocked by hardware-aware reformulations that maximize GPU utilization.

References

- Ian S. Abramson. On bandwidth variation in kernel estimates—a square root law. *The Annals of Statistics*, 10(4):1217–1223, 1982. doi: 10.1214/aos/1176345986.
- Arturs Backurs, Piotr Indyk, and Tal Wagner. Space and time efficient kernel density estimation in high dimensions. In *Advances in Neural Information Processing Systems*, volume 32, 2019.
- Zdravko I. Botev, Joseph F. Grotowski, and Dirk P. Kroese. Kernel density estimation via diffusion. *The Annals of Statistics*, 38(5):2916–2957, 2010. doi: 10.1214/10-AOS799.
- Moses Charikar and Paris Siminelakis. Hashing-based-estimators for kernel density in high dimensions. In *2017 IEEE 58th Annual Symposium on Foundations of Computer Science (FOCS)*, pages 1032–1043, 2017. doi: 10.1109/FOCS.2017.99.
- Moses Charikar and Paris Siminelakis. Multi-resolution hashing for fast pairwise summations. In *2019 IEEE 60th Annual Symposium on Foundations of Computer Science (FOCS)*, pages 769–792, 2019. doi: 10.1109/FOCS.2019.00051.
- Moses Charikar, Michael Kapralov, Navid Nouri, and Paris Siminelakis. Kernel density estimation through density constrained near neighbor search. In *2020 IEEE 61st Annual Symposium on Foundations of Computer Science (FOCS)*, pages 172–183, 2020. doi: 10.1109/FOCS46700.2020.00025.

Benjamin Charlier, Jean Feydy, Joan Alexis Glaunès, François-David Collin, and Ghislain Durif. Kernel operations on the gpu, with autodiff, without memory overflows. *Journal of Machine Learning Research*, 22(74):1–6, 2021. URL <http://jmlr.org/papers/v22/20-275.html>.

Benjamin Coleman and Anshumali Shrivastava. Sub-linear RACE sketches for approximate kernel density estimation on streaming data. In *Proceedings of The Web Conference 2020 (WWW)*, pages 1739–1749, 2020. doi: 10.1145/3366423.3380244.

Tri Dao, Daniel Y. Fu, Stefano Ermon, Atri Rudra, and Christopher Ré. Flashattention: Fast and memory-efficient exact attention with IO-awareness. *arXiv preprint arXiv:2205.14135*, 2022.

Raaz Dwivedi and Lester Mackey. Kernel thinning. In *Proceedings of the 34th Conference on Learning Theory (COLT)*, volume 134, 2021.

Elliot L. Epstein, Rajat Vadiraj Dwaraknath, Thanawat Sornwanee, John Winnicki, and Jerry Weihong Liu. SD-KDE: Score-debiased kernel density estimation. In *The Thirty-ninth Annual Conference on Neural Information Processing Systems*, 2025. URL <https://openreview.net/forum?id=fvho95EtPu>.

Jianqing Fan and Tien-Chung Hu. Bias correction and higher order kernel functions. *Statistics & Probability Letters*, 13(3):235–243, 1992. doi: 10.1016/0167-7152(92)90053-8.

Leslie Greengard and John Strain. The fast gauss transform. *SIAM Journal on Scientific and Statistical Computing*, 12(1):79–94, 1991. doi: 10.1137/0912004.

Jonathan Ho, Ajay Jain, and Pieter Abbeel. Denoising diffusion probabilistic models. *arXiv preprint arXiv:2006.11239*, 2020.

Aapo Hyvärinen. Estimation of non-normalized statistical models by score matching. *Journal of Machine Learning Research*, 6:695–709, 2005.

M. C. Jones and D. F. Signorini. A comparison of higher-order bias kernel density estimators. *Journal of the American Statistical Association*, 92(439):1063–1073, 1997. doi: 10.1080/01621459.1997.10474062.

Praneeth Kacham, Vahab Mirrokni, and Peilin Zhong. PolySketchFormer: Fast transformers via sketching polynomial kernels. In *Proceedings of the 41st International Conference on Machine Learning (ICML)*, 2024.

Yingzhen Li and Richard E. Turner. Gradient estimators for implicit models. *arXiv preprint arXiv:1705.07107*, 2017.

Chen Luo and Anshumali Shrivastava. Arrays of (locality-sensitive) count estimators (ACE): Anomaly detection on the edge. In *Proceedings of the 2018 World Wide Web Conference (WWW)*, pages 1439–1448, 2018. doi: 10.1145/3178876.3186056.

Paulius Micikevicius, Sharan Narang, Jonah Alben, Gregory Diamos, Erich Elsen, David Garcia, Boris Ginsburg, Michael Houston, Oleksii Kuchaiev, Ganesh Venkatesh, and Hao Wu. Mixed precision training. *arXiv preprint arXiv:1710.03740*, 2017.

Emanuel Parzen. On estimation of a probability density function and mode. *The Annals of Mathematical Statistics*, 33(3):1065–1076, 1962. doi: 10.1214/aoms/1177704472.

Adam Paszke, Sam Gross, Francisco Massa, Adam Lerer, James Bradbury, Gregory Chanan, Trevor Killeen, Zeming Lin, Natalia Gimelshein, Luca Antiga, Alban Desmaison, Andreas Köpf, Edward Yang, Zachary DeVito, Martin Raison, Alykhan Tejani, Sasank Chilamkurthy, Benoit Steiner, Lu Fang, Junjie Bai, and Soumith Chintala. Pytorch: An imperative style, high-performance deep learning library. *arXiv preprint arXiv:1912.01703*, 2019.

Fabian Pedregosa, Gaël Varoquaux, Alexandre Gramfort, Vincent Michel, Bertrand Thirion, Olivier Grisel, Mathieu Blondel, Andreas Müller, Joel Nothman, Gilles Louppe, Peter Prettenhofer, Ron Weiss, Vincent Dubourg, Jake VanderPlas, Alexandre Passos, David Cournapeau, Matthieu Brucher, Matthieu Perrot, and Édouard Duchesnay. Scikit-learn: Machine learning in Python. *Journal of Machine Learning Research*, 12:2825–2830, 2011.

- 383 Ali Rahimi and Benjamin Recht. Random features for large-scale kernel machines. In *Advances in*
384 *Neural Information Processing Systems*, 2007.
- 385 Murray Rosenblatt. Remarks on some nonparametric estimates of a density function. *The Annals of*
386 *Mathematical Statistics*, 27(3):832–837, 1956. doi: 10.1214/aoms/1177728190.
- 387 Abhishek Shetty, Raaz Dwivedi, and Lester Mackey. Distribution compression in near-linear time. In
388 *Proceedings of the 10th International Conference on Learning Representations (ICLR)*, 2022.
- 389 B. W. Silverman. *Density Estimation for Statistics and Data Analysis*. Chapman and Hall, 1986. doi:
390 10.1201/9781315140919.
- 391 Paris Siminelakis, Kexin Rong, Peter Bailis, Moses Charikar, and Philip Levis. Rehashing kernel
392 evaluation in high dimensions. In *Proceedings of the 36th International Conference on Machine*
393 *Learning (ICML)*, volume 97, pages 5789–5798, 2019.
- 394 Yang Song and Stefano Ermon. Generative modeling by estimating gradients of the data distribution.
395 *arXiv preprint arXiv:1907.05600*, 2019.
- 396 Yang Song, Jascha Sohl-Dickstein, Diederik P. Kingma, Abhishek Kumar, Stefano Ermon, and Ben
397 Poole. Score-based generative modeling through stochastic differential equations. *arXiv preprint*
398 *arXiv:2011.13456*, 2021. ICLR 2021.
- 399 Philippe Tillet, H. T. Kung, and David Cox. Triton: An intermediate language and compiler for tiled
400 neural network computations. In *Proceedings of the 3rd ACM SIGPLAN International Workshop*
401 *on Machine Learning and Programming Languages*, 2019. doi: 10.1145/3315508.3329973.
- 402 Pascal Vincent. A connection between score matching and denoising autoencoders. *Neural Computa-*
403 *tion*, 23(7):1661–1674, 2011. doi: 10.1162/NECO_a_00142.
- 404 M. P. Wand and M. C. Jones. *Kernel Smoothing*. Chapman and Hall, 1995. doi: 10.1201/b14876.
- 405 Samuel Williams, Andrew Waterman, and David Patterson. Roofline: An insightful visual perfor-
406 mance model for multicore architectures. *Communications of the ACM*, 52(4):65–76, 2009. doi:
407 10.1145/1498765.1498785.
- 408 Changjiang Yang, Ramani Duraiswami, Nail A. Gumerov, and Larry S. Davis. Improved fast
409 gauss transform and efficient kernel density estimation. In *9th IEEE International Conference*
410 *on Computer Vision (ICCV 2003), 14-17 October 2003, Nice, France*, pages 464–471. IEEE
411 Computer Society, 2003. doi: 10.1109/ICCV.2003.1238383. URL [https://doi.org/10.1109/](https://doi.org/10.1109/ICCV.2003.1238383)
412 [ICCV.2003.1238383](https://doi.org/10.1109/ICCV.2003.1238383).
- 413 Amir Zandieh, Insu Han, Majid Daliri, and Amin Karbasi. KDEformer: Accelerating transformers
414 via kernel density estimation. In *Proceedings of the 40th International Conference on Machine*
415 *Learning (ICML)*, volume 202, pages 40605–40623, 2023.

416 A Additional Experiments

417 This appendix reports additional measurements that complement the benchmarks in Section 6. We
418 extend the runtime sweep to stronger exact baselines and to a second GPU, isolate the contributions
419 of streaming and Tensor-Core tiling through controlled ablations, and evaluate Flash-SD-KDE on
420 a real-data anomaly detection benchmark. All measurements use the 16-D pipeline. Unless stated
421 otherwise, runtimes are reported in milliseconds, exclude one-time warmup and JIT compilation
422 costs, and use the 16-D Gaussian-mixture benchmark of Section 6 with $n_{\text{test}} = n_{\text{train}}/8$.

423 A.1 Extended baseline sweep on the RTX A6000

424 To place Flash-SD-KDE against a fuller set of exact baselines, we extend the PyKeOps comparison
425 from a single point to the full n_{train} sweep and add a `torch.compile` variant of the GEMM-based
426 PyTorch baseline. Table 1 reports runtimes on the RTX A6000.

n_{train}	n_{test}	sklearn	Torch	torch.compile	PyKeOps	Flash-SD-KDE
2,048	256	33.0	0.9	6.0	1.9	0.5
4,096	512	126.2	2.1	2.9	2.5	0.4
8,192	1,024	527.6	7.5	5.0	3.6	0.5
16,384	2,048	2,149.2	28.8	11.3	6.7	1.0
32,768	4,096	8,017.0	113.3	34.3	21.7	2.4

Table 1: Average per-call runtime (ms) of exact 16-D SD-KDE implementations on the RTX A6000, lower is better. Columns are scikit-learn Gaussian KDE, the GEMM-based PyTorch SD-KDE, the same PyTorch implementation under `torch.compile`, the PyKeOps LazyTensor SD-KDE, and our Flash-SD-KDE. Data are drawn from the 16-D Gaussian-mixture benchmark of Section 6, with $n_{\text{test}} = n_{\text{train}}/8$. One-time warmup and JIT compilation costs are excluded.

n_{train}	n_{test}	sklearn	Torch	torch.compile	PyKeOps	Flash-SD-KDE
2,048	256	51.60	0.55	65.20	3.29	0.51
4,096	512	270.94	1.78	63.74	4.15	0.62
8,192	1,024	810.39	6.51	54.25	5.51	1.19
16,384	2,048	3,362.95	25.31	59.88	11.19	3.92

Table 2: Average per-call runtime (ms) on a Tesla V100-PCIE-16GB, using the same 16-D Gaussian-mixture benchmark and $n_{\text{test}} = n_{\text{train}}/8$ protocol as Table 1. Lower is better. One-time warmup and JIT costs are excluded.

Flash-SD-KDE is the fastest exact method at every problem size in the sweep. At $n_{\text{train}} = 32,768$, it is roughly $9.0\times$ faster than PyKeOps, $14.3\times$ faster than `torch.compile`, $47.2\times$ faster than the eager PyTorch baseline, and over three orders of magnitude faster than scikit-learn. The PyTorch-based baselines slow down rapidly with n_{train} , while PyKeOps and Flash-SD-KDE both scale gracefully; among these two, Flash-SD-KDE retains a consistent margin throughout the sweep.

A.2 Generalization across GPU families

Section 6 measures performance on the RTX A6000. To check that the gains are not specific to that architecture, we rerun the same sweep on a Tesla V100-PCIE-16GB. The V100 has a different SFU/FP32 ratio and a smaller Tensor-Core throughput than the A6000, so it provides an independent stress test of the implementation. Table 2 reports the result.

The qualitative ordering is preserved: Flash-SD-KDE remains the fastest method at every measured size on the V100, with margins comparable to those seen on the A6000. The eager PyTorch SD-KDE baseline is competitive on the V100 at small n_{train} but degrades rapidly with n_{train} , while `torch.compile` pays a large warm overhead that the exclusion of JIT costs alone does not remove.

A.3 Effect of query batch size

Many KDE workloads operate at small or moderate query batch sizes, in which case the score pass dominates and the cost of batching across queries is a secondary effect. We isolate this by fixing $n_{\text{train}} = 32,768$ and sweeping n_{test} across four orders of magnitude. Table 3 reports the result.

Flash-SD-KDE remains in the 1.91–2.75 ms range across a $4,096\times$ sweep of n_{test} , while every baseline shows a noticeable rise at the largest query batch. This pattern is consistent with the kernel-level analysis in Section 4: the score pass scales as $O(n_{\text{train}}^2)$ and is independent of n_{test} , so when the score dominates (*i.e.* at large n_{train} relative to n_{test}), the end-to-end runtime is essentially flat in n_{test} . Flash-SD-KDE is therefore particularly attractive in workflows that score a small number of queries against a large training set, such as anomaly or out-of-distribution scoring.

A.4 Kernel launch-parameter sweep

The full launch-parameter sweep is moved here from the main text. This preserves the reproducibility detail while keeping Section 6 focused on the dominant design choices and measured ablations.

n_{train}	n_{test}	Torch	torch.compile	PyKeOps	Flash-SD-KDE
32,768	4	100.97	32.44	15.64	1.92
32,768	16	101.00	32.41	15.43	1.91
32,768	64	103.27	33.44	15.48	1.91
32,768	256	104.09	33.68	16.11	1.92
32,768	1,024	106.22	33.69	15.81	1.97
32,768	4,096	116.81	35.97	15.85	2.11
32,768	16,384	158.18	43.14	16.79	2.75

Table 3: Average per-call runtime (ms) on the RTX A6000 as a function of the query batch size n_{test} , with $n_{\text{train}} = 32,768$ held fixed. Lower is better. Data are drawn from the 16-D Gaussian-mixture benchmark.

Method	Runtime (ms)	Peak alloc (MB)
Streamed Flash + Tensor Cores	10.99	16.25
Streamed Flash + No Tensor Cores	49.52	16.25
Full materialization + Tensor Cores	611.63	16,400.50
Full materialization + No Tensor Cores	617.49	16,400.50

Table 4: Runtime and peak extra GPU memory at $n_{\text{train}} = 65,536$, $n_{\text{test}} = 8,192$ on the RTX A6000. “Streamed Flash” is the production Flash-SD-KDE kernel; “Full materialization” explicitly forms the $n_{\text{train}} \times n_{\text{train}}$ and $n_{\text{test}} \times n_{\text{train}}$ kernel matrices in global memory. Peak alloc is reported above the input-tensor footprint and excludes shared workspace.

454 Nsight Systems traces show that roughly 95% of the SD-KDE runtime at $n_{\text{train}} = 32\text{k}$, $n_{\text{test}} = 4\text{k}$
455 is spent computing the empirical score. We therefore concentrated optimization effort on the score
456 kernel, sweeping the launch parameters to raise utilization. Specifically we varied:

- 457 • $BLOCK_M \in \{32, 64, 128, 256\}$,
- 458 • $BLOCK_N \in \{32, 64, 128, 256, 512, 1024\}$,
- 459 • $num_warps \in \{1, 2, 4, 8\}$,
- 460 • $num_stages \in \{1, 2, 4\}$,

461 and selected the combination that minimized runtime. For this workload we found that
462 $BLOCK_M = 64$, $BLOCK_N = 1024$, $num_warps = 2$, and $num_stages = 2$ gave the
463 best overall performance, increasing measured FLOP utilization by more than $2\times$ relative to smaller-
464 tile settings.

465 A.5 Streaming versus full materialization

466 The streaming design described in Section 4 avoids ever materializing the full pairwise kernel
467 matrices. To quantify how much of the speedup is attributable to streaming alone, we compare the
468 streamed Flash-SD-KDE kernel against an explicit “full-materialization” baseline that allocates the
469 $n_{\text{train}} \times n_{\text{train}}$ training kernel matrix and the $n_{\text{test}} \times n_{\text{train}}$ query kernel matrix. Both implementations
470 otherwise share the same numerical formulation. Table 4 reports runtime and peak GPU memory
471 above the resident input-tensor footprint at the largest size we can materialize on a single A6000,
472 $n_{\text{train}} = 65,536$, $n_{\text{test}} = 8,192$.

473 Streaming reduces peak extra GPU memory by roughly $1,000\times$ at this size, from ~ 16 GB to ~ 16
474 MB, and gives a $55.7\times$ runtime speedup over the materialized Tensor-Core variant. The materialized
475 baseline cannot scale meaningfully beyond this size on the A6000, while streaming Flash-SD-KDE
476 continues to operate well under typical GPU memory budgets. Streaming therefore accounts for the
477 bulk of the memory savings reported in Section 6, separately from the Tensor-Core speedup studied
478 in Section A.6.

A.6 Tensor-Core ablation within the streamed kernel

Within the streamed Flash-SD-KDE kernel we additionally compare a Tensor-Core GEMM path against an otherwise identical FP32 path that disables Tensor-Core tiling. This isolates the Tensor-Core contribution from streaming and from the launch-parameter tuning of Section 4. Table 5 sweeps n_{train} from 2,048 to 65,536 with $n_{\text{test}} = n_{\text{train}}/8$.

n_{train}	n_{test}	TC (ms)	no-TC (ms)	no-TC / TC
2,048	256	0.30	0.31	$1.04\times$
4,096	512	0.28	0.35	$1.24\times$
8,192	1,024	0.29	0.87	$2.99\times$
16,384	2,048	0.66	2.83	$4.27\times$
32,768	4,096	2.12	10.48	$4.94\times$
65,536	8,192	9.30	44.59	$4.80\times$

Table 5: Tensor-Core ablation within the streamed Flash-SD-KDE kernel on the RTX A6000. “TC” is the production kernel; “no-TC” is an otherwise identical kernel with Tensor-Core tiling disabled. The right-most column reports $\text{runtime}_{\text{no-TC}}/\text{runtime}_{\text{TC}}$, so values above $1\times$ mean the Tensor-Core path is faster. Lower runtime is better.

The Tensor-Core advantage grows from a marginal $1.04\times$ at the smallest size to $4.94\times$ at $n_{\text{train}} = 32,768$, and stabilizes near $5\times$ at the largest sizes we can run. This is consistent with arithmetic-intensity reasoning: at small n_{train} the GEMM tiles are too small to saturate Tensor Cores, while at moderate-to-large n_{train} the GEMM portion dominates and the kernel approaches the mixed Tensor-Core/FP32 roofline analysed in Section 4. Together with Section A.5, this gives a clean attribution: streaming is responsible for the memory-footprint reduction at all sizes, while Tensor-Core tiling drives the runtime gain at scale.

We do not currently fuse the score pass with the KDE pass into a single kernel. Nsight Systems traces show that the score pass accounts for roughly 90% of the end-to-end runtime, so fusing the two kernels yields only a small incremental gain at the cost of substantial additional code complexity. We retain the simpler two-kernel structure.

A.7 Anomaly detection on the ACE benchmark

The benchmarks above measure runtime and oracle accuracy. To check that Flash-SD-KDE is competitive on a real, non-synthetic task, we evaluate it on the anomaly-detection benchmark of Luo and Shrivastava [2018]. We follow the same protocol: rank all points by an outlier score, threshold at the top- K to recover a candidate outlier set, and report the number of true outliers correctly recovered, the total number reported, the number missed, and the wall-clock execution time. Flash-SD-KDE produces an outlier score by negating its density estimate, so that low-density points rank as outliers.

We compare against the twelve baselines reported in the ACE paper on two of the three datasets evaluated there: the Statlog Shuttle dataset ($n = 34,987$, $d = 9$, 879 outliers) and the ALOI object-image dataset ($n = 50,000$, $d = 27$, 1,508 outliers). The non-Flash-SD-KDE rows in Tables 6–7 are taken directly from the ACE paper; the Flash-SD-KDE rows are our measurements on the RTX A6000 using the same datasets and the same top- K protocol.

On Shuttle, Flash-SD-KDE recovers 807 of the 879 true outliers at the chosen threshold, compared with 610 for the strongest non-Flash baseline (kNNW) and 273 for ACE itself, while running in 0.60 s versus the next fastest method, ACE at 0.81 s. On ALOI, Flash-SD-KDE is competitive on recovery quality (within 20% of the best non-Flash baseline) and is more than an order of magnitude faster than ACE, and roughly three orders of magnitude faster than the distance/density-based baselines. The two datasets together cover relatively low ($d = 9$) and moderate ($d = 27$) dimensionalities, both of which are within the regime where the Tensor-Core formulation of Section 4 is effective.

We additionally evaluate Flash-SD-KDE at large scale on a public point-anomaly release of the KDD-Cup 99 HTTP dataset ($n = 620,098$, $d = 29$, 1,052 labelled anomalies), which is comparable in scale to the 596,853-point KDD HTTP subset used in the ACE paper. On this release Flash-SD-KDE reaches a ROC-AUC of 0.92 in 7.27 s of end-to-end execution time on the RTX A6000. For reference, the ACE paper reports 23.33 s on its own KDD HTTP subset. We do not claim a

Method	Reported	Correct	Missed	Time (s)	vs ACE
ACE	6,763	273	606	0.81	1×
LOF	4,356	381	498	14.12	17.4×
kNN	4,897	493	386	12.35	15.2×
kNNW	5,264	610	269	13.54	16.7×
LoOP	6,145	201	678	14.51	17.9×
LDOF	6,433	330	549	16.42	20.3×
ODIN	9,775	375	504	12.21	15.1×
KDEOS	12,630	314	565	11.73	14.5×
COF	9,133	280	599	13.45	16.6×
LDF	9,809	375	504	19.93	24.6×
INFLO	4,488	183	696	14.03	17.3×
FastVOA	8,532	271	608	235.10	290.2×
Flash-SD-KDE	5,130	807	72	0.60	0.74×

Table 6: Anomaly recovery on Statlog Shuttle ($n = 34,987$, $d = 9$, 879 true outliers). “Reported” is the number of points returned at the chosen threshold, “Correct” the number of those that are true outliers, and “Missed” the number of true outliers not recovered. Time is total execution time in seconds; “vs ACE” is the runtime ratio relative to the ACE row. Non-Flash-SD-KDE rows are reproduced from the ACE paper; Flash-SD-KDE is measured on the RTX A6000.

Method	Reported	Correct	Missed	Time (s)	vs ACE
ACE	7,216	340	1,168	1.26	1×
LOF	4,476	519	989	72.31	57.4×
kNN	5,428	447	1,061	63.27	50.2×
kNNW	5,558	329	1,179	89.96	71.4×
LoOP	5,121	253	1,255	59.97	47.6×
LDOF	7,501	470	1,038	60.39	47.9×
ODIN	10,110	162	1,346	72.69	57.6×
KDEOS	9,515	404	1,104	55.89	44.4×
COF	8,746	284	1,224	81.74	64.9×
LDF	9,133	301	1,207	60.51	48.0×
INFLO	10,328	420	1,088	72.13	57.2×
FastVOA	8,931	319	1,189	291.10	231.0×
Flash-SD-KDE	7,250	437	1,071	0.09	0.07×

Table 7: Anomaly recovery on the ALOI object-image dataset ($n = 50,000$, $d = 27$, 1,508 true outliers). Columns and source attribution match Table 6.

like-for-like comparison here because the public release does not match the exact subset used in the ACE paper, but the result indicates that Flash-SD-KDE remains practical at the $\sim 10^6$ -point scale on commodity hardware.

B Low-dimensional and Laplace-correction details

B.1 Flash-SD-KDE in low dimensions

For completeness we briefly summarize the 1-D SD-KDE arithmetic intensity and empirical behavior. With $n_{\text{train}} = k$ and $n_{\text{test}} = k/8$, there are two steps:

- Score + shift:** For each training point we compute $\hat{s}(x_i)$ from all k points and then form $x_i^{\text{SD}} = x_i + \frac{h^2}{2} \hat{s}(x_i)$. This uses $O(k^2)$ pairwise kernel interactions. With the RTX A6000 hardware ratio (128 FP32 ALUs, 16 SFUs per SM) we budget an exp as 8 flop-equivalents. The score accumulation requires one exp and roughly eight additional arithmetic ops (subtraction, scaling, accumulation), yielding $c_1 \approx 16$ flops per (train, train) pair.
- KDE on debiased samples:** We then evaluate a standard Gaussian KDE at $k/8$ query points using the k debiased samples, which uses $O(k^2/8)$ interactions. Each pair requires one exp (8 flops) plus about six other operations (difference, square, scaling, accumulation), so we take $c_2 \approx 14$ flops per (train, test) pair.

535 The total work is therefore approximated by

$$\text{FLOPs}(k) \approx c_1 k^2 + c_2 k (k/8) \approx 16k^2 + 14 \frac{k^2}{8} = 17.75 k^2.$$

536 For $k = 32k$ this is on the order of 2×10^{10} flops. When we fix $n_{\text{test}} = k/8$ we move approximately
 537 5k bytes (counting one read of each train and test point and one write of each output), so the arithmetic
 538 intensity scales as

$$I(k) = \frac{\text{FLOPs}(k)}{\text{Bytes}(k)} \approx \frac{17.75 k^2}{5 k} \approx 3.55 k \quad \text{flops/byte},$$

539 placing realistic problem sizes firmly in the compute-bound regime.

540 Empirically we sweep k over powers of two from 1024 to 64k, with $n_{\text{test}} = k/8$, and average
 541 over three seeds. For each configuration we record scikit-learn Gaussian KDE time, and SD-KDE
 542 Torch time, and Flash-SD-KDE time. Across the entire range, the Flash-SD-KDE implementation is
 543 consistently faster than scikit-learn, with speedups growing with k ; at the largest sizes, Flash-SD-KDE
 544 achieves around 2 orders-of-magnitude speedup relative to the sklearn baseline, as shown in Figure 4.




Figure 4: Average runtime (log-scale y -axis) of scikit-learn KDE, SD-KDE (Torch), and Flash-SD-KDE across n_{train} in 1-D; annotations show observed runtimes.

545 As shown in Figure 5, we also compare the utilization of Flash-SD-KDE with SD-KDE (Torch)
 546 over a range of 1-D problem sizes, using powers of two for n_{train} up to $2^{16} \approx 64$ thousand (and
 547 $n_{\text{test}} = n_{\text{train}}/8$).




Figure 5: Utilization of the 1-D Flash-SD-KDE kernel and SD-KDE (Torch) for large n_{train} (powers of two up to 2^{16}). Labels show the observed runtime for each data size; error bars are omitted because a single seed is used per point.

548 B.2 Laplace-corrected KDE derivation and 1D benchmarks




Figure 6: Oracle error on a 1D mixture-of-Gaussians benchmark. We report MISE and MIAE versus n_{train} for KDE, Flash-Laplace-KDE (fused Laplace correction), non-fused Laplace correction, and Flash-SD-KDE. The Laplace-corrected estimators can be slightly negative, so error is computed on the signed density.




Figure 7: Runtime and speedup for Laplace correction in 1D. The left panel shows total runtime for the fused Flash-Laplace-KDE kernel and a non-fused implementation. The right panel reports speedup ratios, including Flash-SD-KDE relative to Flash-Laplace-KDE for context.

549 We evaluate Laplace-corrected KDE on a 1D oracle density. Figure 6 compares KDE, Flash-
550 Laplace-KDE (fused Laplace correction), non-fused Laplace correction, and Flash-SD-KDE. The
551 Laplace-corrected estimators can take small negative values for large $\|x - x_i\|$. We therefore compute
552 errors on the signed density, and we report the negative-mass diagnostics separately in the benchmark
553 logs.

554 Figure 7 shows the runtime advantage of fusion. The fused Flash-Laplace-KDE kernel consistently
555 outperforms the non-fused Laplace correction across n_{train} . For context, the same plot also reports the
556 Flash-SD-KDE to Flash-Laplace-KDE runtime ratio.

557 As an intuitive explanation of why both methods achieve the same leading bias reduction, we can first
558 view vanilla KDE² as a heat semigroup (kernel) $P_t = e^{\frac{t}{2}\Delta}$ with $t = h^2$ applying to the empirical
559 density function f_n (sampled with replacement) to be

$$\hat{f} = P_t f_n.$$

²Throughout, we assume that the kernel is Gaussian for simplicity.

560 From the linearity of the heat semigroup, we then have that the expected estimate

$$\mathbb{E} [\hat{f}] = \mathbb{E} [P_t f_n] = P_t \mathbb{E} [f_n] = P_t f,$$

561 where f is the true probability density function.

562 The bias is then, under a smoothness assumption we assume throughout,

$$\mathbb{E} [\hat{f}] - f = (P_t - I)f = \frac{t}{2} \Delta f + O(t^2).$$

563 The Laplace kernel deals with this leading term by using an operator $(I - \frac{t}{2} \Delta) P_t$ instead, making
564 its bias

$$\mathbb{E} [\hat{f}^{\text{LC}}] - f = \left[\left(I - \frac{t}{2} \Delta \right) P_t - I \right] f = O(t^2).$$

565 The SD-KDE algorithm approaches this debias task differently by applying a **non-linear**³ transporta-
566 tion kernel $T_{\frac{t}{2}} \nabla \log f$ before applying a standard heat kernel P_t , making its bias

$$\begin{aligned} \mathbb{E} [\hat{f}^{\text{SD-KDE}}] - f &= \left[P_t T_{\frac{t}{2}} \nabla f - I \right] f \\ &= \left[P_t \left(f - \frac{t}{2} \nabla \cdot (f \nabla \log f) + O(t^2) \right) - f \right] \\ &= \left[P_t \left(f - \frac{t}{2} \Delta f + O(t^2) \right) - f \right] \\ &= \left[P_t \left(I - \frac{t}{2} \Delta \right) - I \right] f + O(t^2) \\ &= O(t^2). \end{aligned}$$

567 Recall that $t = h^2$, so we have that Laplace KDE and SD-KDE have an $O(h^4)$ bias, while vanilla
568 KDE has an $O(h^2)$ bias.

569 Note that the analysis above assumes that the SD-KDE has access to the oracle, so the translation
570 kernel is linear. In practice, the score is empirical, making the translation kernel

$$T_{\frac{t}{2}} \nabla \log(P_{t'} f_n)$$

571 where $t' = \frac{h^2}{2}$ is the bandwidth for the kernel used for score estimation. Thus, the empirical SD-KDE
572 is

$$\hat{f}^{\text{Emp SD-KDE}} = T_{\frac{t}{2}} \nabla \log(P_{t'} f_n) f_n,$$

573 which is no longer linear in f_n .

574 C Additional Literature Review

575 This appendix expands the related-work discussion of Section 2 along three axes that are most
576 relevant to score-debiased and other bias-corrected density estimators: sub-quadratic approximate
577 KDE algorithms, sample compression and coresets methods, and KDE-as-primitive applications in
578 modern machine learning. We also discuss potential hybrids of these algorithmic directions with the
579 hardware-accelerated exact pipeline of Section 4. The discussion is intentionally complementary to
580 the experimental focus of the main paper: the works below approximate either the kernel sum or the
581 training set, while our work makes the exact SD-KDE estimator practical at scale.

³Even when the transportation kernel has a non-linear transportation. The kernel itself is linear.

C.1 Sub-quadratic approximate KDE via locality-sensitive hashing

A line of work starting with the hashing-based estimator (HBE) framework of Charikar and Siminelakis [2017] reduces the cost of evaluating a Gaussian kernel sum from quadratic to sub-quadratic by sampling from locality-sensitive hash buckets. At a high level, HBE hashes both the training set and the query point into LSH buckets whose collision probability is calibrated to the kernel, and uses bucket cardinalities together with a small number of refinement samples to estimate the kernel sum to a target relative error. Backurs et al. [2019] improve the construction so that it requires only linear preprocessing time and space while retaining the same query time, and Siminelakis et al. [2019] bring the scheme closer to practice through adaptive sample-size selection and sharper data-dependent variance bounds. Charikar and Siminelakis [2019] extend the construction beyond Gaussian-like kernels to a much broader class of pairwise summations using a multi-resolution hashing scheme. A subsequent refinement [Charikar et al., 2020] casts the sub-quadratic KDE problem as a density-constrained near-neighbor search and removes the redundancy that arises in HBE when nearby points fall into different buckets on opposite sides of the query.

The ACE method of Luo and Shrivastava [2018], which we evaluate against in Section A.7, is a practical instantiation of this family specialized for streaming anomaly detection: rather than estimating the kernel sum, it tracks per-bucket counts and uses them as proxies for local density. The closely related RACE sketch of Coleman and Shrivastava [2020] extends the LSH-counting view to streaming KDE with explicit error guarantees. ACE was a natural choice for our anomaly-detection comparison precisely because it shares the high-level inverse-density-as-anomaly viewpoint of a KDE-based scorer while bringing the sub-quadratic LSH cost.

These methods are complementary to ours rather than substitutes. HBE-family algorithms approximate *classical* KDE, so applying them out of the box to SD-KDE would forfeit the bias-reduction guarantees that motivate SD-KDE in the first place. A natural hybrid direction (Section C.5) is to use a sub-quadratic approximation only for the empirical-score pass and retain exact streaming Tensor-Core evaluation for the final density.

C.2 Sample compression, coresets, and kernel thinning

Compression methods change the statistical estimator by replacing the sample, whereas the experiments in this paper evaluate exact SD-KDE acceleration at fixed n ; we therefore treat them as complementary future-work context rather than direct exact-runtime baselines.

A second line of work reduces the cost of kernel computations by replacing the full training set with a much smaller weighted subset. Dwivedi and Mackey [2021] introduce kernel thinning, which produces a coreset of size $O(\sqrt{n})$ whose worst-case kernel-mean error matches that of the full sample up to logarithmic factors, by halving the dataset under a self-balancing process and selecting the half that better preserves a target reproducing-kernel Hilbert space functional. Shetty et al. [2022] extend this construction with a meta-procedure (Compress++) that reduces the total runtime to near-linear while inflating the integration error by at most a constant factor, making kernel thinning practical at substantially larger scales.

Compression-based methods affect a different cost axis from the one we target. They reduce n before any kernel evaluations are performed, while we keep n fixed and accelerate the resulting $O(n^2)$ computation. For SD-KDE specifically, kernel thinning could in principle be used to reduce both the training set and the empirical-score sample, at the cost of inflating bias and variance by amounts that have not yet been characterized for score-debiased estimators. We view a careful theoretical and empirical study of kernel thinning under SD-KDE as out of scope for this paper but a promising follow-up.

C.3 Higher-order bias correction in KDE

The Laplace-corrected estimator of Section 5 belongs to the broader family of higher-order bias kernels. Classical work in this area includes the higher-order kernel constructions of Fan and Hu [1992] and the systematic comparison of Jones and Signorini [1997], both already discussed in the main related-work section. The novelty of our use of this construction is not in the kernel itself but in identifying it as the right first-order surrogate for SD-KDE in the limit where the empirical score is

replaced by its leading-order Taylor expansion, and in giving a fused GPU realization that retains the same effective bias reduction at a fraction of the runtime cost.

C.4 KDE as a primitive in modern ML systems

Although the present paper targets density estimation directly, KDE-style sums also arise as computational primitives inside modern deep-learning pipelines. Self-attention with a softmax kernel can be written as a normalized weighted sum that is closely related to a Gaussian KDE evaluated at the query, and several recent works exploit this view to design sub-quadratic attention approximations. Zandieh et al. [2023] treat the attention denominator as a KDE and use a sub-quadratic KDE algorithm in the spirit of the FOCS line above to approximate it, obtaining provable spectral-norm guarantees and large end-to-end speedups. Kacham et al. [2024] replace the softmax kernel with a polynomial kernel and use polynomial-kernel sketches from randomized numerical linear algebra to obtain a linear-time attention mechanism with approximation guarantees. Streaming-style LSH sketches in the spirit of RACE [Coleman and Shrivastava, 2020] have similarly been adapted as the scoring primitive inside attention modules. FlashAttention [Dao et al., 2022] is the closest analogue on the hardware-aware side: it does not approximate the kernel sum but reorders the exact computation to fit the GPU memory hierarchy and to expose Tensor-Core GEMMs, which is the same high-level recipe we apply to SD-KDE.

These works underline that fast KDE-style primitives have downstream value beyond density estimation. The implementation pattern in this paper (matrix-multiplication re-ordering, streaming accumulation, mixed Tensor-Core/FP32 work) is a natural building block for any future architecture in which a score-corrected or higher-order kernel sum is used in place of vanilla attention or a vanilla KDE primitive.

C.5 Hybrid algorithmic–hardware directions for future work

The literature surveyed above suggests several hybrid directions that pair algorithmic reductions with the hardware-accelerated exact pipeline of this paper.

Approximate-score, exact-evaluation hybrids. The empirical score in SD-KDE is itself a sum of weighted kernels, and the SD-KDE estimator is empirically robust to small score perturbations. This makes it natural to plug a sub-quadratic approximate KDE algorithm, such as those of Charikar and Siminelakis [2017], Backurs et al. [2019], Siminelakis et al. [2019], Charikar and Siminelakis [2019], Charikar et al. [2020], into the score-evaluation step while retaining the exact Tensor-Core kernel of Section 4 for the final density evaluation. Quantifying how approximate-score error propagates into the bias-reduction guarantee of SD-KDE is an open theoretical question in this direction.

Spectral and grid-based score acceleration. For low-to-moderate dimensions, the score can also be computed by evaluating the gradient of a smoothed density on a grid, with non-uniform fast Fourier transforms (NUFFT) reducing the cost of the underlying sum. Combining a NUFFT-based score with the exact Flash-SD-KDE evaluation kernel is a promising hybrid for the dimensionalities we target, and is on our roadmap.

Coreset-then-accelerate. A complementary hybrid uses kernel thinning [Dwivedi and Mackey, 2021, Shetty et al., 2022] or another coreset construction to reduce n before invoking the exact pipeline, trading a controlled error increase for a square-root reduction in compute. Because Flash-SD-KDE is already efficient at large n , the regime where this trade is attractive likely lies at the very largest scales, where even a hardware-accelerated exact $O(n^2)$ pass becomes prohibitive.

In all three directions, a fast and accurate exact SD-KDE primitive is useful as the inner loop and as a ground-truth reference for evaluating approximation error, which is the role this paper aims to fill.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [Yes]

Justification: The abstract and introduction state the main contributions: a Tensor-Core reordering, a streaming Triton implementation, and the Flash-Laplace-KDE surrogate. The reported claims are supported by the method and experiments in Sections 4, 5, and 6 and the appendix.

Guidelines:

- The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: The paper discusses the main scope limits, including the signed nature of the Laplace correction, and future work on broader bandwidths, kernels, and multi-GPU support. These points appear mainly in Sections 5 and 7.

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [N/A]

Justification: N/A. The paper is mainly a systems and empirical contribution. It includes derivations and references prior SD-KDE theory, but it does not introduce new formal theoretical results that require a full assumptions-and-proofs section.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: The paper gives the algorithmic formulation, hardware details, launch parameters, baselines, datasets, and runtime tables needed to reproduce the main results. It also provides a code repository link.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).

- (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: The paper links the Flash-SD-KDE code repository and describes the synthetic and public benchmark settings used in the experiments. The appendix gives additional baselines and ablations for reproducing the main claims.

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [Yes]

Justification: The experimental setup is described at a practical level: hardware, dimensionality, sample and query sizes, baselines, launch-parameter sweeps, and benchmark datasets are all specified in Sections 3 and 6 and the appendix.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: The paper reports error bars suitably, such as in figures 2, 3, and 4.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: The paper reports the main compute environment, including the RTX A6000 workstation and the V100 comparison, and gives measured runtimes for the experiments. Section 3 and the appendix provide the most relevant compute details.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: The work uses synthetic and public benchmark data and does not involve human subjects, sensitive data collection, or release of high-risk models. Thus, the paper conforms in every respect with the NeurIPS Code of Ethics.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes]

Justification: The Impact Statement notes the main positive impact, namely making SD-KDE practical at larger scales, and states that the work does not introduce unique societal risks beyond standard density-estimation use cases.

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]

Justification: The paper does not release a high-risk model, scraped dataset, or system such as a language model or image generator. The released asset is an implementation of a density-estimation method.

Guidelines:

- The answer [N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [Yes]

Justification: The paper cites the main existing software libraries, baselines, and datasets used in the experiments, including PyTorch, scikit-learn, Triton, and the public anomaly-detection benchmarks. The software dependencies used for implementation and baselines are open-source packages, including BSD-licensed PyTorch and scikit-learn, and are used only as experimental dependencies or comparison points rather than redistributed as part of a new asset.

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [Yes]

Justification: The main new asset is the Flash-SD-KDE implementation. The paper links the repository and documents the core method, hardware assumptions, benchmarks, and ablations in the paper and appendix.

Guidelines:

- The answer [N/A] means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [N/A]

Justification: The paper does not involve crowdsourcing, user studies, or other research with human subjects.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.

990 • According to the NeurIPS Code of Ethics, workers involved in data collection, curation,
991 or other labor should be paid at least the minimum wage in the country of the data
992 collector.

993 **15. Institutional review board (IRB) approvals or equivalent for research with human**
994 **subjects**

995 Question: Does the paper describe potential risks incurred by study participants, whether
996 such risks were disclosed to the subjects, and whether Institutional Review Board (IRB)
997 approvals (or an equivalent approval/review based on the requirements of your country or
998 institution) were obtained?

999 Answer: [N/A]

1000 Justification: The paper does not involve crowdsourcing or research with human subjects, so
1001 IRB approval or equivalent review is not applicable.

1002 Guidelines:

1003 • The answer [N/A] means that the paper does not involve crowdsourcing nor research
1004 with human subjects.

1005 • Depending on the country in which research is conducted, IRB approval (or equivalent)
1006 may be required for any human subjects research. If you obtained IRB approval, you
1007 should clearly state this in the paper.

1008 • We recognize that the procedures for this may vary significantly between institutions
1009 and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the
1010 guidelines for their institution.

1011 • For initial submissions, do not include any information that would break anonymity (if
1012 applicable), such as the institution conducting the review.

1013 **16. Declaration of LLM usage**

1014 Question: Does the paper describe the usage of LLMs if it is an important, original, or
1015 non-standard component of the core methods in this research? Note that if the LLM is used
1016 only for writing, editing, or formatting purposes and does *not* impact the core methodology,
1017 scientific rigor, or originality of the research, declaration is not required.

1018 Answer: [N/A]

1019 Justification: The core method does not use LLMs as an important, original, or non-standard
1020 component.

1021 Guidelines:

1022 • The answer [N/A] means that the core method development in this research does not
1023 involve LLMs as any important, original, or non-standard components.

1024 • Please refer to our LLM policy in the NeurIPS handbook for what should or should not
1025 be described.